

# Dimensionality Reduction and Hypothesis Testing

Jawadul Chowdhury

April 10, 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Abstract . . . . .	3
1.2	About the Dataset . . . . .	3
<b>2</b>	<b>Method</b>	<b>3</b>
2.1	Feature Extraction . . . . .	3
2.2	Dimensionality Reduction . . . . .	3
2.3	The Clustering Process . . . . .	3
2.3.1	Centroid Based Clustering . . . . .	3
2.3.2	Agglomerative Hierarchical Clustering . . . . .	3
2.4	Document Frequencies of Words . . . . .	3
2.5	Enriched Words with Statistical Testing . . . . .	3
2.5.1	Using Binomial Testing . . . . .	3
2.5.2	The Top 200 Words . . . . .	4
<b>3</b>	<b>Results</b>	<b>4</b>
3.1	The Explained Variance Ratio . . . . .	4
3.2	Scatter Plot of Two Components . . . . .	4
3.3	Colored Scatter Plots of Two Components . . . . .	4
3.4	The Silhouette Score . . . . .	5
3.5	Centroid Based Clustering . . . . .	5
3.6	Agglomerative Hierarchical Clustering . . . . .	5
3.7	The Confusion Matrix Heatmap . . . . .	5
<b>4</b>	<b>Discussion</b>	<b>6</b>
4.1	Understanding the Dataset . . . . .	6
4.2	The Feature Matrix . . . . .	6
4.2.1	Rows and Columns of the Feature Matrix . . . . .	6
4.2.2	Calculating the Sparsity Ratio . . . . .	6
4.2.3	The Meaning Behind The Feature Matrix . . . . .	6
4.2.4	The Compressed Row Format . . . . .	6
4.3	A Closer Look at Dimensionality Reduction . . . . .	6
4.4	A Closer Look at Clustering . . . . .	6
4.4.1	Cluster Assignment Analysis . . . . .	6
4.4.2	Choice of Clustering Algorithm . . . . .	6
4.4.3	Application of Transformations . . . . .	7
4.4.4	The Ham & Spam Labels vs Cluster Labels . . . . .	7
4.5	The Top 200 Words . . . . .	7

# 1 Introduction

In this section, we discuss about what the paper is about as well as the kind of dataset we use, as well as the features and how big the dataset is.

## 1.1 Abstract

In this paper, we analyze 63,542 emails. We convert the raw text from these emails into a feature matrix using a "bag of words" model. Each column of the feature matrix corresponds to one word, each row corresponds to one email, and the entry stores the number of times that word was found in the email. We then perform dimensionality reduction, cluster the emails in two clusters. Lastly, we perform binomial testing on all words in each cluster and filter out the top 200 words.

## 1.2 About the Dataset

Using the 63,452 emails, we converted them into a pandas data frame. A closer look at the pandas data frame tells us that it consists of 5 features, which are `category`, `to.address`, `from.address`, `subject` and `body`. The dataset consists of 63542 rows and 5 columns. For the experiments we perform on the dataset, we are primarily concerned with the `category` and `body`, as will be seen in the sections below.

# 2 Method

In this section, we discuss the methods we used to perform dimensionality reduction and feature extraction. We discuss in detail all the steps that were performed and why they were performed. The purpose of this section is detail the steps and explain the concepts behind what we did.

## 2.1 Feature Extraction

Throughout the course of this lab, we extract all the emails that we have stored in a directory from a .json format into a pandas data frame. We then perform feature extraction, where we create a feature matrix from the text that makes up the body of the email. We ensure that we only capture words that are mentioned more than 10 times. The reason we do this is because we would like to reduce noise from rare words as well as prevent overfitting in the unsupervised learning model that will be used later on.

Another point worth noting is that when we reduce the number of words, it helps to lower the dimensionality of the data. Lastly, if we were to include a great number of rare words, this will cause inconsistencies and will disrupt the clustering process when using a unsupervised machine learning model.

## 2.2 Dimensionality Reduction

With the current feature matrix that we've created, we perform principal component analysis by reducing the number of dimensions in the data. By reducing it to 10 columns, we have data that is easier to analyze. Next,

we obtain the explained variance ratios, which tell us how much variance is retained by each component after performing dimensionality reduction.

## 2.3 The Clustering Process

After performing dimensionality reduction, we use unsupervised learning models to perform clustering. The two methods of clustering we use are centroid based clustering and agglomerative hierarchical clustering.

### 2.3.1 Centroid Based Clustering

When it comes to centroid based clustering, we used k-means. K-means works by partitioning a set of data points into K clusters based on their features so that data points within the same cluster are more similar to each other than to those in other clusters. Using k-means means we need to determine the value of k (number of clusters), and to find the correct value for k, we determine this using the silhouette score.

### 2.3.2 Agglomerative Hierarchical Clustering

Agglomerative Hierarchical Clustering works by building a hierarchy of clusters. It works by starting each data point as its own individual cluster and then merging the closest cluster at each step, progressively forming larger clusters until all data points belong to one cluster. For the linkage methods, we used single linkage, which finds the distance between two clusters in the shortest distance between any two points from different clusters.

## 2.4 Document Frequencies of Words

After clustering all the emails, we now analyze the clusters we've created and how the words we've captured play a role in clustering. We achieve this by creating a separate matrix for each cluster containing the rows for the points in that cluster. We convert these matrices into a CSC format due to the benefit that it is optimized for column slicing. Next, we calculate the document frequency of each word in each cluster. We perform document clustering as it allows us to determine the importance of a word within a cluster and analyze what the cluster represents.

## 2.5 Enriched Words with Statistical Testing

The aim here is to find words that are enriched in each cluster. As a result, we can interpret the themes in the cluster using statistical tests.

### 2.5.1 Using Binomial Testing

We use binomial testing to know whether a word appears in more emails than expected in one cluster compared to another. The null and alternative hypothesis for binomial testing are states as follows:

- **Null Hypothesis:** the relative document frequencies of the observed cluster are less or equal to those of the tested

- **Alternative Hypothesis:** the document frequency is higher in cluster 0 than in cluster 1

At the end of the day, we use statistical testing to avoid bias.

### 2.5.2 The Top 200 Words

We perform binomial testing on all words in each cluster, and then filter out the top 200 words based on their p-value, as the p-value tells us how strong the evidence is, and whether to null hypothesis is true or if there isn't enough evidence to back the null hypothesis.

## 3 Results

In this section, we look at the visualizations that have been produced, and discuss these visualization in detail that will better explain the decisions made behind choosing the k-value, which unsupervised learning model to use and many more.

### 3.1 The Explained Variance Ratio

In a nutshell, the explained variance ratio tells us how much information each principal component captures from the original data. It shows how important each component is in representing the original data. We plot the explained variance ratios of the components which are derived from reducing the dimensionality of the feature matrix as shown in Figure 1.




Figure 1: Bar Chart of Explained Variance Ratio

When we look at figure 1, it can be seen that the two highest explained variance ratios can be found in component 0 and component 1.

### 3.2 Scatter Plot of Two Components

Moving on, using the two components with the highest explained variance ratios, we now create a scatter plot where the x-axis is component 1 and the y-axis is component 2 as shown in figure 2 below.




Figure 2: Scatter Plot of Two Components

When we look at the scatter plot above, it can be clearly seen that there are two clusters, or maybe two groupings. Lets forget the clusters for a minute. Lets pay attention to the axes. When we plot component 1 against component 2, what we really are doing is that we're visualizing the data in a two dimensional area using two components that were extracted via dimensionality reduction. The components are explained as follows:

- **Component 1:** it is the direction in the feature space that accounts for the most variance
- **Component 2:** it is orthogonal to component 1 and captures the second most variance

### 3.3 Colored Scatter Plots of Two Components

After creating the first scatter plot, we now create a second scatter plot. Using the `category` feature which consists of just two unique values, which are ham and spam, we now color the plots on the scatter plot according to whether they are ham and spam as shown in Figure 3.




Figure 3: Scatter Plot of (Ham & Spam)

When we look at our new scatter plot above, we can see that the ham and spam clusters are intermingled together, which is not quite helpful as we would like to have two separate clusters. The intermingling of these

two components can be a indication that it is these two components that aren't enough to distinguish the categories of ham and spam.

### 3.4 The Silhoutte Score

Before we have an attempt at making a k-means model, we need to determine the value of k in order to determine the number of clusters we'd like to create. The silhouette score is a metric used to evaluate the quality of a clustering result. It allows us to understand how well each data point fits within its assigned cluster compared to other clusters. We create a line graph of the silhouette scores against the values of k as shown in Figure 4 below.




Figure 4: Line Graph of Silhouette Scores

When we look at the line graph, it can be evidently seen that the best k-value would be 2, as it has the highest silhouette score. As a result, we use this k-value when it comes to centroid based and agglomerative hierarchial clustering.

### 3.5 Centroid Based Clustering

The first unsupervised learning model we use is k-means. Using a k-value of 2, we cluster using the two SVD components. The clustering algorithm labels the points so that all points in the same cluster have the same cluster ID. We color the points according to their cluster labels as shown in Figure 5 below.




Figure 5: K-Means Clustering

When we look at figure 5, we can tell that the clustering has been done wrong because there is significant overlap between the two clusters, especially around the center of the plot. As a result, we turn to agglomerative hierarchial clustering.

### 3.6 Agglomerative Hierarchial Clustering

The next unsupervised learning model we use is agglomerative hierarchial clustering. Setting the number of clusters as 2 and the type of linkage as single, we extract the top two components and produce Figure 6.

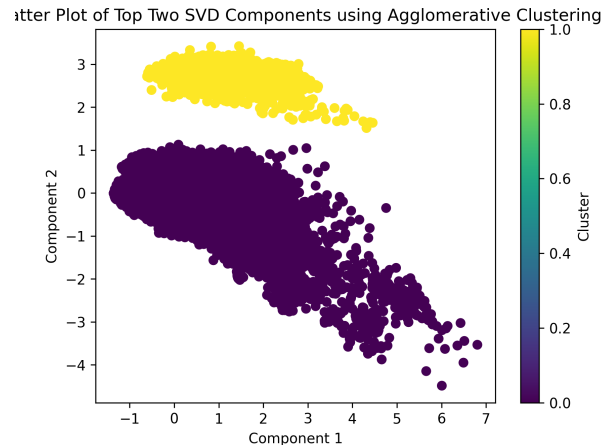


Figure 6: Agglomerative Hierarchial Clustering

When we look at figure 6, the clustering here looks much better. This is because there is clear separation between the clusters and agglomerative clustering is much better suited towards non-convex clusters.

### 3.7 The Confusion Matrix Heatmap

Using the results of the agglomerative hierarchial clustering model, we at last generate a confusion matrix to compare the ham / spam labels to the cluster labels using the confusion matrix that we have generate as shown in Figure 7.




Figure 7: HeatMap of Confusion Matrix

The results of the heat maps shows us that the number of true positives (emails correctly identified as spam) is 5793. The number of true negatives (emails correctly identified as ham) is 40450. When it came to false positives (ham incorrectly identified as spam) we got 0 and false negatives (spam incorrectly identified as ham) is 17,299.

Zero false positives implies that every positive results identified by the clustering model was incorrect, which is a great result to have. However, a few false negatives means that the model misses actual positives at times. We can now use the results from the confusion matrix to calculate the F1 score, which measures how well the clusters match the actual spam vs ham classification, as it balances both precision and recall as shown in equation 1.

$$\text{F1 Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (1)$$

When we calculate the F1 score, we get 0.42 when rounded to two decimal places. Since our F1 score is close to 0.5, the clusters match the categories partially.

## 4 Discussion

In this last section, we discuss the conclusions from the results we've obtained and summarize what we've understood from the experiments that we've conducted.

### 4.1 Understanding the Dataset

When we first converted the json files into a pandas data frame, we wanted to look at the features that made up the data frame. The names of the features were `category`, `to_address`, `from_address`, `subject` & `body`. When we use `.info()` to know about the data types of the features, they all appear to be objects, however, there is more to that. The feature `category` is categorical as there are only two unique values, which makes it exception. The rest of the features are objects because there are simply bodies of text.

### 4.2 The Feature Matrix

Now, we look at the feature matrix that we produced and look at the key statistics, such as the number of rows and columns, non-zero entries and the sparsity ratio. We also explain the structure and meaning of the matrix, and provide an estimate of the memory usage.

#### 4.2.1 Rows and Columns of the Feature Matrix

When we create a feature matrix by using `fit_transform()` from `CountVectorizer`, we can see that the matrix consists of a total of 63542 rows and 32144 columns. When we look at the number on non-zero entries in the matrix, we get 6388795 non-zero entries. Non-zero entries refers to the positions in the matrix where the value is not zero. It is simply counting how many times a word appears in any email.

#### 4.2.2 Calculating the Sparsity Ratio

When it comes to calculating the sparsity ratio, we calculate it by  $100 \times (\text{number of non-zero entries} / \text{maximum possible entries})$ . Using equation 2 below, we get the sparsity ratio as 0.3127938123616995.

$$\text{Sparsity} = 100 \times \left( \frac{\text{Non-zero entries}}{\text{Total entries}} \right) \quad (2)$$

#### 4.2.3 The Meaning Behind The Feature Matrix

When we look at the feature, the number of rows are the number of emails, from 0 to 63541 emails. The number of columns are the words in the emails that appear atleast more than 10 emails. We then cross-check these words against the emails to see if the word is present in the email. If the word is present the email, we use a binary value of 1 to indicate it as true, else 0 to indicate it as false.

#### 4.2.4 The Compressed Row Format

Lastly, the vectorize we use returns a sparse matrix in CSR. Since we assume that the sparse matrix uses one 32-bit and one 32-bit integer for each non-zero entry and one 32-bit integer for each row. When we calculate the memory usage as shown in equation 3, we get the memory as near 49 MB.

$$\text{Total Memory Usage} = 8 \times \text{nnz} + 4 \times (m + 1) \quad (3)$$

### 4.3 A Closer Look at Dimensionality Reduction

When it came to performing dimensionality reduction, we used `fit_transform` to create a new matrix with 10 components. In order to find the components with the largest explained variance ratios, we found that it was component 0 and 1 which had a explained variance ratio of 0.057 and 0.028 respectively.

### 4.4 A Closer Look at Clustering

In this section, we look at the clusters we've created using two unsupervised learning models in great detail and discuss what we did.

#### 4.4.1 Cluster Assignment Analysis

Looking at the clustering algorithm that we used (agglomerative hierarchical clustering), we observe a convex structure to these clusters. It is evident that the data points naturally form two distinct clusters. When we compare the clusters with the actual labels, we believe that clustering will allow for recovering the label as the clustering forms two distinct groups.

#### 4.4.2 Choice of Clustering Algorithm

When it came k-means clustering, it appears to show irregular cluster shapes. This led to overlap between the two clusters, and where one cluster is bigger than the other. When it came to using agglomerative hierarchical clustering, there is a clear separation between

